

< sudo >

Sudo-ku

Energieeffektivgruppen

October 31, 2016



Planned Operation

Design Overview

Argumentation

Data flow and memory

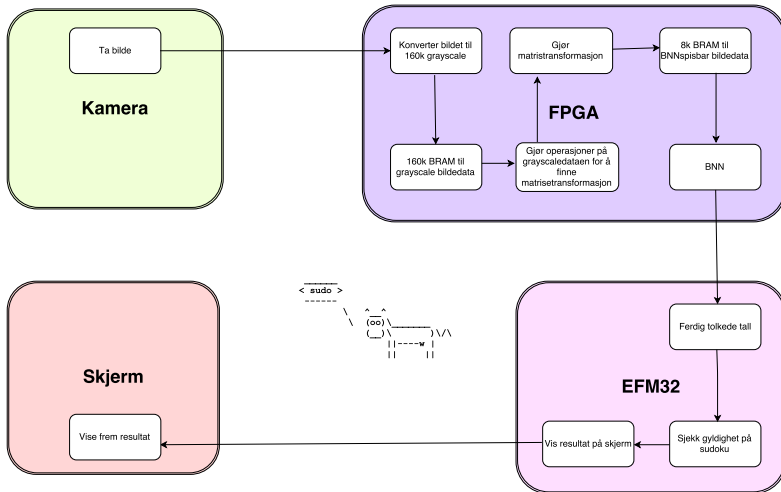
Image transformation

Binarizing the neural network



- ▶ Point camera at an image of a solved sudoku-board
- ▶ Push button
- ▶ Camera image is stored in FPGA BRAM
- ▶ FPGA figures out the transform required to rotate, scale and crop the input image
- ▶ FPGA performs transform, cutting it into pieces, trimming edges of the squares to avoid sending outlines into the BNN etc.
- ▶ BNN (on FPGA) detects digits
- ▶ Digits sent to MCU, where the sudoku is checked and result (and potentially other relevant info, incorrect rows etc.) is output to the user

Planned operation of sudo-ku



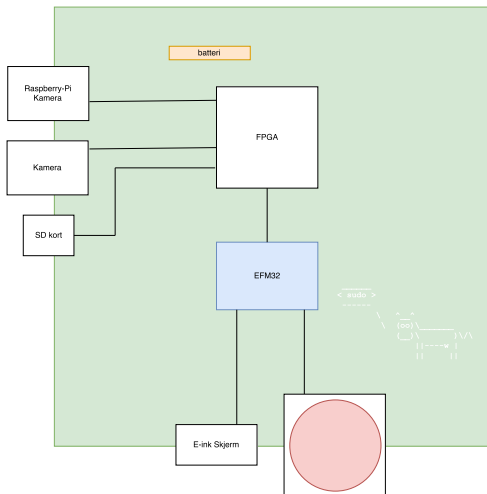
Planned operation of sudo-ku

Picture requirements



- ▶ All but the board itself should be black, the board itself white
- ▶ The camera must be pointed roughly orthogonally to the board; we do not do any perspective transformations
- ▶ The camera can be rotated along the camera-pointing axis by 45 degree either way. This requirement might go away if we find a simple way to figure out the orientation without it.

Design Overview





- ▶ We need 9x9 times 28x28 pixels of relevant image data. Results in $9 \times 9 \times 28 \times 28 \times 2 = 127008$ pixels.
- ▶ Smallest standard accommodating this is HVGA ($480 \times 320 = 153600$ pixels). But is uncommonly supported. Hence we use VGA ($640 \times 480 = 307200$ pixels).
- ▶ We use 1 bit black/white to store the images. Thus we need 307200 bits of storage.
- ▶ Image from the camera is in a different format than our 1-bit black/white. This needs conversion on the fly.
- ▶ Considering the size of the image, it is possible to do keep the image in the memory of the FPGA without the use of an external SRAM



- ▶ With the restrictions we put on the picture we can find the corners with a somewhat naive algorithm that should be possible to implement on the FPGA
- ▶ The process of doing an affine transformation on the camera input data is a simple, arithmetical operation, easy to implement elegantly on an FPGA.
- ▶ The resulting, transformed image/images are only needed on the FPGA



- ▶ Output, which includes communicating with a display, is a lot more difficult implementing in hardware, with no real benefit.
- ▶ The MCU is used to control the operation flow of *sudo-ku*.



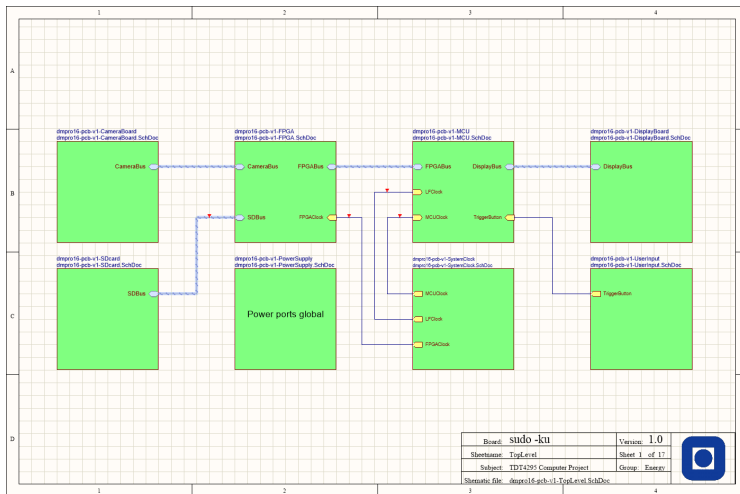
- ▶ Super easy to do on MCU.
- ▶ Have already made a first draft of about 50 lines of code checking a sudoku.
- ▶ Lots of eye-candy can be implemented when everything works
- ▶ Only 700 ways to solve a sudoku. The entire board can be transferred from the FPGA to the MCU in 10 bits. Data transferred is negligible.



- ▶ In case the camera processing does not work, we will still have a way to demo the working BNN by loading pre-formatted, square image data, skipping the transformation steps.


PCB Design

Top level schematic



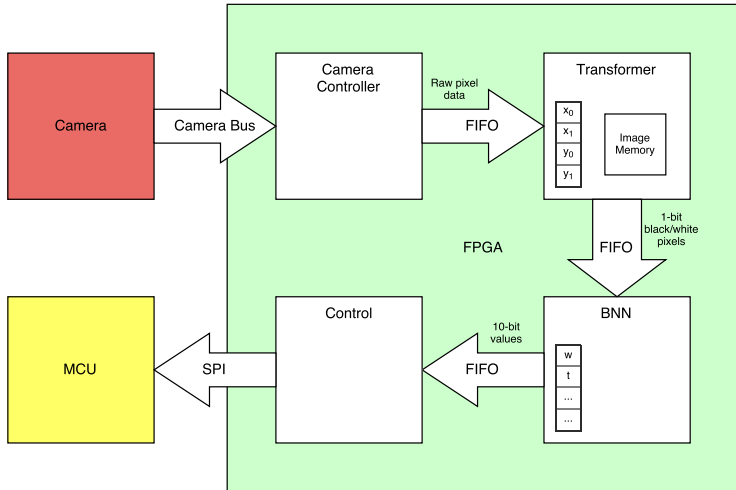
PCB Design

3D Drawing



Data flow and memory

On the FPGA



Data flow and memory

In the Transformer-module

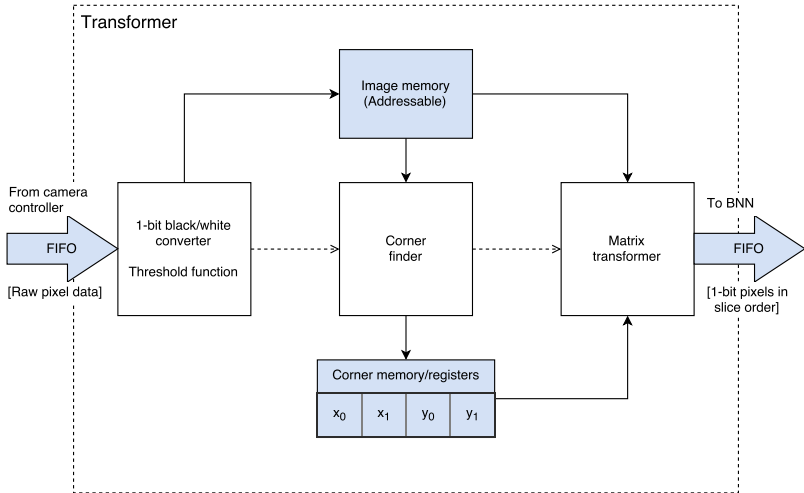



Image transformation

Cropping and straightening



- ▶ Done by a matrix transform *from* the required destination image coordinates (pixels) *to* input-image coordinates (sub-pixel)
- ▶ Uses coordinates of the topmost corners



- ▶ Done by a matrix transform *from* the required destination image coordinates (pixels) *to* input-image coordinates (sub-pixel)
- ▶ Uses coordinates of the topmost corners

$$T(x, y) = \begin{bmatrix} a & -b \\ b & a \end{bmatrix} \cdot \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} x_0 \\ y_0 \end{bmatrix}$$

$$a = \frac{x_1 - x_0}{w_{\text{dest}}}$$

$$b = \frac{y_1 - y_0}{w_{\text{dest}}}$$

Image transformation

Slicing and padding




16

- ▶ A *slice* is a small square containing one digit
- ▶ The pixels are to be sent into the BNN slicewise

- ▶ A *slice* is a small square containing one digit
- ▶ The pixels are to be sent into the BNN slicewise
- ▶ Going from slice coordinates (n, x', y') , n being the slice number, to sudoku-wide coordinates (x, y) with padding w_p :

$$x = (28 + 2w_p)(n \bmod 9) + x' + w_p$$

$$y = (28 + 2w_p)\lfloor n/9 \rfloor + y' + w_p$$



- ▶ A *slice* is a small square containing one digit
- ▶ The pixels are to be sent into the BNN slicewise
- ▶ Going from slice coordinates (n, x', y') , n being the slice number, to sudoku-wide coordinates (x, y) with padding w_p :

$$x = (28 + 2w_p)(n \bmod 9) + x' + w_p$$

$$y = (28 + 2w_p)\lfloor n/9 \rfloor + y' + w_p$$

- ▶ This makes $w_{\text{dest}} = 28 \cdot 9 + 2 \cdot 9 \cdot w_p = 252 + 18w_p$

Binarizing the neural network

Changing to -1 and 1



- ▶ We can represent most of the neural network as -1 or 1 without much accuracy loss
- ▶ The input 0-255 can be -1 if ≤ 127 else 1
- ▶ New weights become $\text{sign}(\text{old weights})$
- ▶ The activation function $\tanh(x)$ can also be replaced by $\text{sign}(x)$
- ▶ Now only the batch normalization is not represented by -1 and 1 yet

Binarizing the neural network

From multiplication to XNOR



	-1	1
1	-1	1
-1	1	-1

Multiplication

	0	1
1	0	1
0	1	0

XNOR

Binarizing the neural network

Finding the batch normalization threshold



- ▶ $y = \gamma \sigma^{-1}(x - \bar{x}) + \beta$
- ▶ $y = \gamma \sigma^{-1}(2x - \bar{x} - prev) + \beta$
- ▶ Find where the output is 0, so we can use only one value
- ▶ $x = 0.5(\bar{x} - \frac{\beta}{\gamma \sigma^{-1}} + prev)$

Binarizing the neural network

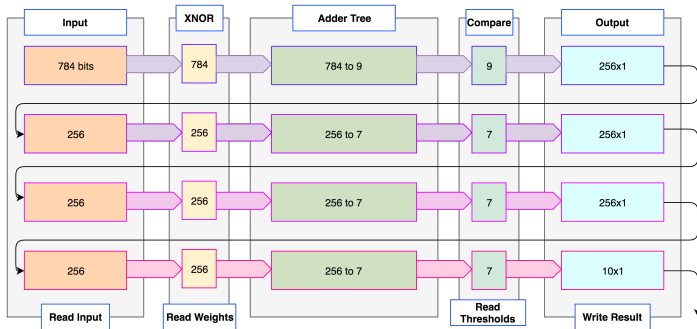
The binarized result



- ▶ $\sum w \otimes out$
- ▶ 0 if the sum is below the batchnorm threshold
- ▶ 1 otherwise

BNN Architecture

Image of the architecture





- ▶ The next neuron in one layer can start processing one clock cycle after the current neuron
- ▶ We can start on the next image as soon as the previous one is done with the first layer
- ▶ All layers can be used at the same time, each working on a different image